Biostat 203B Homework 4

Due Mar 9 @ 11:59PM

Khoa Vu 705600710

Display machine information:

```r
sessionInfo()
```

```
R version 4.4.2 (2024-10-31)
Platform: x86_64-pc-linux-gnu
Running under: Ubuntu 24.04.1 LTS

Matrix products: default
BLAS:   /usr/lib/x86_64-linux-gnu/blas/libblas.so.3.12.0
LAPACK: /usr/lib/x86_64-linux-gnu/lapack/liblapack.so.3.12.0

locale:
 [1] LC_CTYPE=C.UTF-8       LC_NUMERIC=C           LC_TIME=C.UTF-8
 [4] LC_COLLATE=C.UTF-8     LC_MONETARY=C.UTF-8    LC_MESSAGES=C.UTF-8
 [7] LC_PAPER=C.UTF-8       LC_NAME=C              LC_ADDRESS=C
[10] LC_TELEPHONE=C         LC_MEASUREMENT=C.UTF-8 LC_IDENTIFICATION=C

time zone: America/Los_Angeles
tzcode source: system (glibc)

attached base packages:
[1] stats     graphics  grDevices utils     datasets  methods   base

loaded via a namespace (and not attached):
 [1] compiler_4.4.2    fastmap_1.2.0     cli_3.6.4         tools_4.4.2
 [5] htmltools_0.5.8.1 rstudioapi_0.17.1 yaml_2.3.10       rmarkdown_2.29
 [9] knitr_1.49        jsonlite_1.9.1    xfun_0.51         digest_0.6.37
[13] rlang_1.1.5       evaluate_1.0.3
```

Display my machine memory.

```
memuse::Sys.meminfo()
```

```
Totalram:  11.686 GiB
Freeram:    8.145 GiB
```

Load database libraries and the tidyverse frontend:

```
library(bigrquery)
library(dbplyr)
library(DBI)
library(gt)
library(gtsummary)
library(tidyverse)
```

```
-- Attaching core tidyverse packages ---------------------- tidyverse 2.0.0 --
v dplyr     1.1.4     v readr     2.1.5
v forcats   1.0.0     v stringr   1.5.1
v ggplot2   3.5.1     v tibble    3.2.1
v lubridate 1.9.4     v tidyr     1.3.1
v purrr     1.0.4
-- Conflicts ------------------------------------------ tidyverse_conflicts() --
x dplyr::filter() masks stats::filter()
x dplyr::ident()  masks dbplyr::ident()
x dplyr::lag()    masks stats::lag()
x dplyr::sql()    masks dbplyr::sql()
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to becom
```

Q1. Compile the ICU cohort in HW3 from the Google BigQuery database

Below is an outline of steps. In this homework, we exclusively work with the BigQuery database and should not use any MIMIC data files stored on our local computer. Transform data as much as possible in BigQuery database and `collect()` the tibble **only at the end of Q1.7**.

Q1.1 Connect to BigQuery

Authenticate with BigQuery using the service account token. Please place the service account token (shared via BruinLearn) in the working directory (same folder as your qmd file). Do **not** ever add this token to your Git repository. If you do so, you will lose 50 points.

```
# path to the service account token
satoken <- "biostat-203b-2025-winter-4e58ec6e5579.json"
# BigQuery authentication using service account
bq_auth(path = satoken)
```

Connect to BigQuery database `mimiciv_3_1` in GCP (Google Cloud Platform), using the project billing account `biostat-203b-2025-winter`.

```
# connect to the BigQuery database `biostat-203b-2025-mimiciv_3_1`
con_bq <- dbConnect(
    bigrquery::bigquery(),
    project = "biostat-203b-2025-winter",
    dataset = "mimiciv_3_1",
    billing = "biostat-203b-2025-winter"
)
con_bq
```

```
<BigQueryConnection>
  Dataset: biostat-203b-2025-winter.mimiciv_3_1
  Billing: biostat-203b-2025-winter
```

List all tables in the `mimiciv_3_1` database.

```
dbListTables(con_bq)
```

```
 [1] "admissions"        "caregiver"        "chartevents"
 [4] "d_hcpcs"           "d_icd_diagnoses"  "d_icd_procedures"
 [7] "d_items"           "d_labitems"       "datetimeevents"
[10] "diagnoses_icd"     "drgcodes"         "emar"
[13] "emar_detail"       "hcpcsevents"      "icustays"
[16] "ingredientevents"  "inputevents"      "labevents"
[19] "microbiologyevents" "omr"             "outputevents"
[22] "patients"          "pharmacy"         "poe"
[25] "poe_detail"        "prescriptions"    "procedureevents"
[28] "procedures_icd"    "provider"         "services"
[31] "transfers"
```

Q1.2 `icustays` data

Connect to the `icustays` table.

```
# full ICU stays table
icustays_tble <- tbl(con_bq, "icustays") |>
  arrange(subject_id, hadm_id, stay_id) |>
  # show_query() |>
  print(width = Inf)
```

```
# Source:     SQL [?? x 8]
# Database:   BigQueryConnection
# Ordered by: subject_id, hadm_id, stay_id
   subject_id  hadm_id  stay_id first_careunit
        <int>    <int>    <int> <chr>
 1   10000032 29079034 39553978 Medical Intensive Care Unit (MICU)
 2   10000690 25860671 37081114 Medical Intensive Care Unit (MICU)
 3   10000980 26913865 39765666 Medical Intensive Care Unit (MICU)
 4   10001217 24597018 37067082 Surgical Intensive Care Unit (SICU)
 5   10001217 27703517 34592300 Surgical Intensive Care Unit (SICU)
 6   10001725 25563031 31205490 Medical/Surgical Intensive Care Unit (MICU/SICU)
 7   10001843 26133978 39698942 Medical/Surgical Intensive Care Unit (MICU/SICU)
 8   10001884 26184834 37510196 Medical Intensive Care Unit (MICU)
 9   10002013 23581541 39060235 Cardiac Vascular Intensive Care Unit (CVICU)
10   10002114 27793700 34672098 Coronary Care Unit (CCU)
   last_careunit                                   intime
   <chr>                                           <dttm>
 1 Medical Intensive Care Unit (MICU)              2180-07-23 14:00:00
 2 Medical Intensive Care Unit (MICU)              2150-11-02 19:37:00
 3 Medical Intensive Care Unit (MICU)              2189-06-27 08:42:00
 4 Surgical Intensive Care Unit (SICU)             2157-11-20 19:18:02
 5 Surgical Intensive Care Unit (SICU)             2157-12-19 15:42:24
 6 Medical/Surgical Intensive Care Unit (MICU/SICU) 2110-04-11 15:52:22
 7 Medical/Surgical Intensive Care Unit (MICU/SICU) 2134-12-05 18:50:03
 8 Medical Intensive Care Unit (MICU)              2131-01-11 04:20:05
 9 Cardiac Vascular Intensive Care Unit (CVICU)    2160-05-18 10:00:53
10 Coronary Care Unit (CCU)                        2162-02-17 23:30:00
   outtime             los
   <dttm>             <dbl>
 1 2180-07-23 23:50:47 0.410
 2 2150-11-06 17:03:17 3.89
 3 2189-06-27 20:38:27 0.498
```

```
 4 2157-11-21 22:08:00 1.12
 5 2157-12-20 14:27:41 0.948
 6 2110-04-12 23:59:56 1.34
 7 2134-12-06 14:38:26 0.825
 8 2131-01-20 08:27:30 9.17
 9 2160-05-19 17:33:33 1.31
10 2162-02-20 21:16:27 2.91
# i more rows
```

Q1.3 `admissions` data

Connect to the `admissions` table.

```
# # TODO
admissions_tble <- tbl(con_bq, "admissions") |>
  arrange(subject_id, hadm_id) |>
  # show_query() |>
  print(width = Inf)
```

```
# Source:     SQL [?? x 16]
# Database:   BigQueryConnection
# Ordered by: subject_id, hadm_id
   subject_id  hadm_id admittime           dischtime           deathtime
        <int>    <int> <dttm>              <dttm>              <dttm>
 1   10000032 22595853 2180-05-06 22:23:00 2180-05-07 17:15:00 NA
 2   10000032 22841357 2180-06-26 18:27:00 2180-06-27 18:49:00 NA
 3   10000032 25742920 2180-08-05 23:44:00 2180-08-07 17:50:00 NA
 4   10000032 29079034 2180-07-23 12:35:00 2180-07-25 17:55:00 NA
 5   10000068 25022803 2160-03-03 23:16:00 2160-03-04 06:26:00 NA
 6   10000084 23052089 2160-11-21 01:56:00 2160-11-25 14:52:00 NA
 7   10000084 29888819 2160-12-28 05:11:00 2160-12-28 16:07:00 NA
 8   10000108 27250926 2163-09-27 23:17:00 2163-09-28 09:04:00 NA
 9   10000117 22927623 2181-11-15 02:05:00 2181-11-15 14:52:00 NA
10   10000117 27988844 2183-09-18 18:10:00 2183-09-21 16:30:00 NA
   admission_type     admit_provider_id admission_location     discharge_location
   <chr>              <chr>             <chr>                  <chr>
 1 URGENT             P49AFC            TRANSFER FROM HOSPITAL HOME
 2 EW EMER.           P784FA            EMERGENCY ROOM         HOME
 3 EW EMER.           P19UTS            EMERGENCY ROOM         HOSPICE
 4 EW EMER.           P06OTX            EMERGENCY ROOM         HOME
 5 EU OBSERVATION     P39NWO            EMERGENCY ROOM         <NA>
```

```
 6 EW EMER.         P42H7G          WALK-IN/SELF REFERRAL  HOME HEALTH CARE
 7 EU OBSERVATION   P35NE4          PHYSICIAN REFERRAL     <NA>
 8 EU OBSERVATION   P40JML          EMERGENCY ROOM         <NA>
 9 EU OBSERVATION   P47EY8          EMERGENCY ROOM         <NA>
10 OBSERVATION ADMIT P13ACE         WALK-IN/SELF REFERRAL  HOME HEALTH CARE
   insurance language marital_status race  edregtime
   <chr>     <chr>    <chr>          <chr> <dttm>
 1 Medicaid  English  WIDOWED        WHITE 2180-05-06 19:17:00
 2 Medicaid  English  WIDOWED        WHITE 2180-06-26 15:54:00
 3 Medicaid  English  WIDOWED        WHITE 2180-08-05 20:58:00
 4 Medicaid  English  WIDOWED        WHITE 2180-07-23 05:54:00
 5 <NA>      English  SINGLE         WHITE 2160-03-03 21:55:00
 6 Medicare  English  MARRIED        WHITE 2160-11-20 20:36:00
 7 Medicare  English  MARRIED        WHITE 2160-12-27 18:32:00
 8 <NA>      English  SINGLE         WHITE 2163-09-27 16:18:00
 9 Medicaid  English  DIVORCED       WHITE 2181-11-14 21:51:00
10 Medicaid  English  DIVORCED       WHITE 2183-09-18 08:41:00
   edouttime              hospital_expire_flag
   <dttm>                              <int>
 1 2180-05-06 23:30:00                     0
 2 2180-06-26 21:31:00                     0
 3 2180-08-06 01:44:00                     0
 4 2180-07-23 14:00:00                     0
 5 2160-03-04 06:26:00                     0
 6 2160-11-21 03:20:00                     0
 7 2160-12-28 16:07:00                     0
 8 2163-09-28 09:04:00                     0
 9 2181-11-15 09:57:00                     0
10 2183-09-18 20:20:00                     0
# i more rows
```

Q1.4 `patients` data

Connect to the `patients` table.

```
# # TODO
patients_tble <- tbl(con_bq, "patients") |>
  arrange(subject_id) |>
  # show_query() |>
  print(width = Inf)
```

```
# Source:     SQL [?? x 6]
```

```
# Database:   BigQueryConnection
# Ordered by: subject_id
   subject_id gender anchor_age anchor_year anchor_year_group dod
        <int> <chr>       <int>       <int> <chr>             <date>
 1   10000032 F              52        2180 2014 - 2016       2180-09-09
 2   10000048 F              23        2126 2008 - 2010       NA
 3   10000058 F              33        2168 2020 - 2022       NA
 4   10000068 F              19        2160 2008 - 2010       NA
 5   10000084 M              72        2160 2017 - 2019       2161-02-13
 6   10000102 F              27        2136 2008 - 2010       NA
 7   10000108 M              25        2163 2014 - 2016       NA
 8   10000115 M              24        2154 2017 - 2019       NA
 9   10000117 F              48        2174 2008 - 2010       NA
10   10000161 M              60        2163 2020 - 2022       NA
# i more rows
```

Q1.5 `labevents` data

Connect to the `labevents` table and retrieve a subset that only contain subjects who appear in `icustays_tble` and the lab items listed in HW3. Only keep the last lab measurements (by `storetime`) before the ICU stay and pivot lab items to become variables/columns. Write all steps in *one* chain of pipes.

```
# # TODO
subset_itemid <- c(50912, 50971, 50983, 50902, 50882, 51221, 51301, 50931)

#Connecting to d_labitems to get lab item names
d_labitems_tble <- tbl(con_bq, "d_labitems") |>
  arrange(itemid) |>
  # show_query() |>
  print(width = Inf)
```

```
# Source:     SQL [?? x 4]
# Database:   BigQueryConnection
# Ordered by: itemid
   itemid label                              fluid category
    <int> <chr>                              <chr> <chr>
 1  50801 Alveolar-arterial Gradient         Blood Blood Gas
 2  50802 Base Excess                        Blood Blood Gas
 3  50803 Calculated Bicarbonate, Whole Blood Blood Blood Gas
 4  50804 Calculated Total CO2               Blood Blood Gas
```

```
 5  50805 Carboxyhemoglobin                Blood Blood Gas
 6  50806 Chloride, Whole Blood            Blood Blood Gas
 7  50808 Free Calcium                     Blood Blood Gas
 8  50809 Glucose                          Blood Blood Gas
 9  50810 Hematocrit, Calculated           Blood Blood Gas
10  50811 Hemoglobin                       Blood Blood Gas
# i more rows
```

```r
labevents_tble <- tbl(con_bq, "labevents") |>
  #Filtering by measurements we want to take
  filter(itemid %in% subset_itemid) |>
  #We keep only the columns we need
  select(subject_id, itemid, storetime, valuenum) |>
  #Mutating columns for proper joins
  #Mutate subject_id and hadm_id as double to join with icustays_tble
  mutate(subject_id = as.double(subject_id)) |>
  #We inner join with icu stays, keeping only the patients with an icu stay
  #We inner join to discard nonmatching rows from both tibbles
  inner_join(icustays_tble, by = c("subject_id"), copy = TRUE) |>
  filter(storetime < intime) |>
  #We group by stay_id and itemid to get the values for each measurement
  #of each stay
  group_by(subject_id, stay_id, itemid) |>
  #We order by the store time, taking the last measurement before intime
  #Using slice_max(), we take a slice of size one for the highest time
  slice_max(order_by = storetime, n = 1) |>
  summarize(valuenum = mean(valuenum, na.rm = TRUE)) |>
  #We ungroup to get the dataframe back to a normal size
  ungroup() |>
  #We join labevents_tble with d_labitems by itemid
  left_join(d_labitems_tble, by = "itemid", copy = TRUE) |>
  #We subset labevents_tbl to only the columns we need in our final result
  select(c(subject_id, stay_id, valuenum, label)) |>
  #Apply lower case to all labels
  mutate(label = tolower(label)) |>
  #We widen the dataframe to get each row as a subject and ICU stay
  pivot_wider(names_from = label, values_from = valuenum) |>
  #Sorting the tble by subject_id and_stay id for grading purposes
  arrange(subject_id, stay_id) |>
  #Changing white blood cells to wbc, removing spaces
  rename(wbc = `white blood cells`) |>
  # show_query() |>
```

```
print(width = Inf)
```

`summarise()` has grouped output by "subject_id" and "stay_id". You can
override using the `.groups` argument.

Warning: ORDER BY is ignored in subqueries without LIMIT
i Do you need to move arrange() later in the pipeline or use window_order() instead?
ORDER BY is ignored in subqueries without LIMIT
i Do you need to move arrange() later in the pipeline or use window_order() instead?

`summarise()` has grouped output by "subject_id" and "stay_id". You can
override using the `.groups` argument.

Warning: ORDER BY is ignored in subqueries without LIMIT
i Do you need to move arrange() later in the pipeline or use window_order() instead?
ORDER BY is ignored in subqueries without LIMIT
i Do you need to move arrange() later in the pipeline or use window_order() instead?


```
# Source:     SQL [?? x 10]
# Database:   BigQueryConnection
# Ordered by: subject_id, stay_id
   subject_id  stay_id hematocrit sodium bicarbonate chloride   wbc creatinine
        <dbl>    <int>      <dbl>  <dbl>       <dbl>    <dbl> <dbl>      <dbl>
 1   10000032 39553978       41.1    126          25       95   6.9        0.7
 2   10000690 37081114       36.1    137          26      100   7.1        1
 3   10000980 39765666       27.3    144          21      109   5.3        2.3
 4   10001217 34592300       37.4    142          30      104   5.4        0.5
 5   10001217 37067082       38.1    142          22      108  15.7        0.6
 6   10001725 31205490         NA    139          NA       98    NA         NA
 7   10001843 39698942       31.4    138          28       97  10.4        1.3
 8   10001884 37510196       39.7    130          30       88  12.2        1.1
 9   10002013 39060235       34.9    137          24      102   7.2        0.9
10   10002114 34672098       34.3    125          18       NA  16.8        3.1
   glucose potassium
     <dbl>     <dbl>
 1     102       6.7
 2      85       4.8
 3      89       3.9
 4      87       4.1
 5     112       4.2
```

```
 6     NA      4.1
 7    131      3.9
 8    141      4.5
 9    288      3.5
10     95      6.5
# i more rows
```

Q1.6 `chartevents` data

Connect to **chartevents** table and retrieve a subset that only contain subjects who appear
in `icustays_tble` and the chart events listed in HW3. Only keep the first chart events (by
`storetime`) during ICU stay and pivot chart events to become variables/columns. Write all
steps in *one* chain of pipes. Similarly to HW3, if a vital has multiple measurements at the first
`storetime`, average them.

```r
# # TODO
#Taking the subset of item_ids needed
subset_itemid <-c(220045, 220180, 220179, 223761, 220210)

#Connecting to d_items to get item names
d_items_tble <- tbl(con_bq, "d_items") |>
  arrange(itemid) |>
  # show_query() |>
  print(width = Inf)
```

```
# Source:      SQL [?? x 9]
# Database:    BigQueryConnection
# Ordered by: itemid
   itemid label                                abbreviation       linksto
    <int> <chr>                                <chr>              <chr>
 1 220001 Problem List                         Problem List       chartevents
 2 220003 ICU Admission date                   ICU Admission date datetimeevents
 3 220045 Heart Rate                           HR                 chartevents
 4 220046 Heart rate Alarm - High              HR Alarm - High    chartevents
 5 220047 Heart Rate Alarm - Low               HR Alarm - Low     chartevents
 6 220048 Heart Rhythm                         Heart Rhythm       chartevents
 7 220050 Arterial Blood Pressure systolic     ABPs               chartevents
 8 220051 Arterial Blood Pressure diastolic    ABPd               chartevents
 9 220052 Arterial Blood Pressure mean         ABPm               chartevents
10 220056 Arterial Blood Pressure Alarm - Low  ABP Alarm - Low    chartevents
   category             unitname param_type   lownormalvalue highnormalvalue
```

```
   <chr>               <chr>    <chr>                <dbl>        <dbl>
 1 General             <NA>     Text                    NA           NA
 2 ADT                 <NA>     Date and time           NA           NA
 3 Routine Vital Signs bpm      Numeric                 NA           NA
 4 Alarms              bpm      Numeric                 NA           NA
 5 Alarms              bpm      Numeric                 NA           NA
 6 Routine Vital Signs <NA>     Text                    NA           NA
 7 Routine Vital Signs mmHg     Numeric                 90          140
 8 Routine Vital Signs mmHg     Numeric                 60           90
 9 Routine Vital Signs mmHg     Numeric                 NA           NA
10 Alarms              mmHg     Numeric                 NA           NA
# i more rows
```

```r
chartevents_tble <- tbl(con_bq, "chartevents") |>
  #Subset based on needed measurements
  filter(itemid %in% subset_itemid) |>
  #We keep only the columns we need
  select(subject_id, stay_id, itemid, storetime, valuenum) |>
  #Mutate subject_id and stay_id as double to join with icustays_tble
  mutate(subject_id = as.double(subject_id)) |>
  mutate(stay_id = as.double(stay_id)) |>
  #Inner join with ICU Stays
  #Keep patients with an ICU stay
  inner_join(icustays_tble, by = c("subject_id", "stay_id"), copy = TRUE) |>
  #Filter by measurements within the ICU stay
  filter((storetime > intime) & (storetime < outtime)) |>
  #We group by stay_id and itemid to get the values for each measurement
  #of each stay
  group_by(subject_id, stay_id, itemid) |>
  #We order by the store time, taking the first measurement during the ICU stay
  #Using slice_min(), we take a slice of size one for the smallest time in that
  #interval
  slice_min(order_by = storetime, n = 1) |>
  summarize(valuenum = mean(valuenum, na.rm = TRUE)) |>
  #We ungroup to get the dataframe back to a normal size
  ungroup() %>%
  left_join(d_items_tble, by = "itemid", copy = TRUE) |>
  select(c(subject_id, stay_id, valuenum, label)) |>
  #Apply lower case to all labels and remove sapce
  #Note, we have to use the SQL equivalent of str_replace_all
  #Which is REGEXP_REPLACE()
  mutate(label = REGEXP_REPLACE(tolower(label), " ", "_")) |>
```

```
#We widen the dataframe to get each row as a subject and ICU stay
pivot_wider(names_from = label, values_from = valuenum) |>
#Sorting the tble by subject_id and_stay id for grading purposes
arrange(subject_id, stay_id) |>
# show_query() |>
print(width = Inf)
```

`summarise()` has grouped output by "subject_id" and "stay_id". You can
override using the `.groups` argument.

Warning: ORDER BY is ignored in subqueries without LIMIT
i Do you need to move arrange() later in the pipeline or use window_order() instead?
ORDER BY is ignored in subqueries without LIMIT
i Do you need to move arrange() later in the pipeline or use window_order() instead?

`summarise()` has grouped output by "subject_id" and "stay_id". You can
override using the `.groups` argument.

Warning: ORDER BY is ignored in subqueries without LIMIT
i Do you need to move arrange() later in the pipeline or use window_order() instead?
ORDER BY is ignored in subqueries without LIMIT
i Do you need to move arrange() later in the pipeline or use window_order() instead?


```
# Source:     SQL [?? x 7]
# Database:   BigQueryConnection
# Ordered by: subject_id, stay_id
   subject_id  stay_id temperature_fahrenheit respiratory_rate
        <dbl>    <dbl>                  <dbl>            <dbl>
 1   10000032 39553978                   98.7               24
 2   10000690 37081114                   97.7               24.3
 3   10000980 39765666                   98                 23.5
 4   10001217 34592300                   97.6               14
 5   10001217 37067082                   98.5               18
 6   10001725 31205490                   97.7               19
 7   10001843 39698942                   97.9               16.5
 8   10001884 37510196                   98.1               13
 9   10002013 39060235                   97.2               14
10   10002114 34672098                   97.9               21
   non_invasive_blood_pressure_diastolic non_invasive_blood_pressure_systolic
                                   <dbl>                                <dbl>
```

```
 1                                      48                                        84
 2                                    56.5                                       106
 3                                     102                                       154
 4                                    93.3                                       156
 5                                      90                                       151
 6                                      56                                        73
 7                                      78                                       110
 8                                    30.5                                      174.
 9                                      62                                      98.5
10                                      80                                       112
   heart_rate
        <dbl>
 1          91
 2          78
 3          76
 4        79.3
 5          86
 6          86
 7        124.
 8          49
 9          80
10        110.
# i more rows
```

Q1.7 Put things together

This step is similar to Q7 of HW3. Using *one* chain of pipes `|>` to perform following data wrangling steps: (i) start with the `icustays_tble`, (ii) merge in admissions and patients tables, (iii) keep adults only (age at ICU intime $>= 18$), (iv) merge in the labevents and chartevents tables, (v) `collect` the tibble, (vi) sort `subject_id`, `hadm_id`, `stay_id` and `print(width = Inf)`.

```
# # TODO
mimic_icu_cohort <- icustays_tble |>
  #Left join with admissions_tble by subject_id and hadm_id
  left_join(admissions_tble, by = c("subject_id", "hadm_id")) |>
  #Left join with patients_tbl by subject_id
  left_join(patients_tble, by = c("subject_id")) |>
  #Left join with labevents_tble by subject_id and stay_id
  left_join(labevents_tble, by = c("subject_id", "stay_id")) |>
  #Left join with chartevents_tble by subject_id and stay_id
  left_join(chartevents_tble, by = c("subject_id", "stay_id")) |>
```

```
#Creating age at intime %>%
mutate(age_intime = anchor_age + (year(intime) - anchor_year)) |>
#Filter by adults (age_intime >= 18)
filter(age_intime >= 18) |>
#Collecting the tibble
collect() |>
#Sorting the tble by subject_id and_stay id for grading purposes
arrange(subject_id, stay_id) |>
print(width = Inf)
```

`summarise()` has grouped output by "subject_id" and "stay_id". You can
override using the `.groups` argument.
`summarise()` has grouped output by "subject_id" and "stay_id". You can
override using the `.groups` argument.

Warning: ORDER BY is ignored in subqueries without LIMIT
i Do you need to move arrange() later in the pipeline or use window_order() instead?
ORDER BY is ignored in subqueries without LIMIT
i Do you need to move arrange() later in the pipeline or use window_order() instead?
ORDER BY is ignored in subqueries without LIMIT
i Do you need to move arrange() later in the pipeline or use window_order() instead?
ORDER BY is ignored in subqueries without LIMIT
i Do you need to move arrange() later in the pipeline or use window_order() instead?
ORDER BY is ignored in subqueries without LIMIT
i Do you need to move arrange() later in the pipeline or use window_order() instead?
ORDER BY is ignored in subqueries without LIMIT
i Do you need to move arrange() later in the pipeline or use window_order() instead?
ORDER BY is ignored in subqueries without LIMIT
i Do you need to move arrange() later in the pipeline or use window_order() instead?
ORDER BY is ignored in subqueries without LIMIT
i Do you need to move arrange() later in the pipeline or use window_order() instead?
ORDER BY is ignored in subqueries without LIMIT
i Do you need to move arrange() later in the pipeline or use window_order() instead?
ORDER BY is ignored in subqueries without LIMIT
i Do you need to move arrange() later in the pipeline or use window_order() instead?

A tibble: 94,458 x 41
    subject_id hadm_id stay_id first_careunit
         <int>   <int>   <int> <chr>
  1   10000032 29079034 39553978 Medical Intensive Care Unit (MICU)
  2   10000690 25860671 37081114 Medical Intensive Care Unit (MICU)
  3   10000980 26913865 39765666 Medical Intensive Care Unit (MICU)
  4   10001217 27703517 34592300 Surgical Intensive Care Unit (SICU)

```
 5    10001217 24597018 37067082 Surgical Intensive Care Unit (SICU)
 6    10001725 25563031 31205490 Medical/Surgical Intensive Care Unit (MICU/SICU)
 7    10001843 26133978 39698942 Medical/Surgical Intensive Care Unit (MICU/SICU)
 8    10001884 26184834 37510196 Medical Intensive Care Unit (MICU)
 9    10002013 23581541 39060235 Cardiac Vascular Intensive Care Unit (CVICU)
10    10002114 27793700 34672098 Coronary Care Unit (CCU)
   last_careunit                                      intime
   <chr>                                              <dttm>
 1 Medical Intensive Care Unit (MICU)                 2180-07-23 14:00:00
 2 Medical Intensive Care Unit (MICU)                 2150-11-02 19:37:00
 3 Medical Intensive Care Unit (MICU)                 2189-06-27 08:42:00
 4 Surgical Intensive Care Unit (SICU)                2157-12-19 15:42:24
 5 Surgical Intensive Care Unit (SICU)                2157-11-20 19:18:02
 6 Medical/Surgical Intensive Care Unit (MICU/SICU)   2110-04-11 15:52:22
 7 Medical/Surgical Intensive Care Unit (MICU/SICU)   2134-12-05 18:50:03
 8 Medical Intensive Care Unit (MICU)                 2131-01-11 04:20:05
 9 Cardiac Vascular Intensive Care Unit (CVICU)       2160-05-18 10:00:53
10 Coronary Care Unit (CCU)                           2162-02-17 23:30:00
   outtime               los admittime           dischtime
   <dttm>              <dbl> <dttm>               <dttm>
 1 2180-07-23 23:50:47 0.410 2180-07-23 12:35:00 2180-07-25 17:55:00
 2 2150-11-06 17:03:17 3.89  2150-11-02 18:02:00 2150-11-12 13:45:00
 3 2189-06-27 20:38:27 0.498 2189-06-27 07:38:00 2189-07-03 03:00:00
 4 2157-12-20 14:27:41 0.948 2157-12-18 16:58:00 2157-12-24 14:55:00
 5 2157-11-21 22:08:00 1.12  2157-11-18 22:56:00 2157-11-25 18:00:00
 6 2110-04-12 23:59:56 1.34  2110-04-11 15:08:00 2110-04-14 15:00:00
 7 2134-12-06 14:38:26 0.825 2134-12-05 00:10:00 2134-12-06 12:54:00
 8 2131-01-20 08:27:30 9.17  2131-01-07 20:39:00 2131-01-20 05:15:00
 9 2160-05-19 17:33:33 1.31  2160-05-18 07:45:00 2160-05-23 13:30:00
10 2162-02-20 21:16:27 2.91  2162-02-17 22:32:00 2162-03-04 15:16:00
   deathtime           admission_type           admit_provider_id
   <dttm>              <chr>                     <chr>
 1 NA                  EW EMER.                  P06OTX
 2 NA                  EW EMER.                  P26QQ4
 3 NA                  EW EMER.                  P06OTX
 4 NA                  DIRECT EMER.              P276OU
 5 NA                  EW EMER.                  P3610N
 6 NA                  EW EMER.                  P32W56
 7 2134-12-06 12:54:00 URGENT                    P67ATB
 8 2131-01-20 05:15:00 OBSERVATION ADMIT         P49AFC
 9 NA                  SURGICAL SAME DAY ADMISSION P8286C
10 NA                  OBSERVATION ADMIT         P46834
   admission_location     discharge_location insurance language marital_status
```

```
   <chr>                 <chr>               <chr>     <chr>    <chr>
 1 EMERGENCY ROOM        HOME                Medicaid  English  WIDOWED
 2 EMERGENCY ROOM        REHAB               Medicare  English  WIDOWED
 3 EMERGENCY ROOM        HOME HEALTH CARE    Medicare  English  MARRIED
 4 PHYSICIAN REFERRAL    HOME HEALTH CARE    Private   Other    MARRIED
 5 EMERGENCY ROOM        HOME HEALTH CARE    Private   Other    MARRIED
 6 PACU                  HOME                Private   English  MARRIED
 7 TRANSFER FROM HOSPITAL DIED               Medicare  English  SINGLE
 8 EMERGENCY ROOM        DIED                Medicare  English  MARRIED
 9 PHYSICIAN REFERRAL    HOME HEALTH CARE    Medicare  English  SINGLE
10 PHYSICIAN REFERRAL    HOME HEALTH CARE    Medicaid  English  <NA>
   race                  edregtime           edouttime
   <chr>                 <dttm>              <dttm>
 1 WHITE                 2180-07-23 05:54:00 2180-07-23 14:00:00
 2 WHITE                 2150-11-02 11:41:00 2150-11-02 19:37:00
 3 BLACK/AFRICAN AMERICAN 2189-06-27 06:25:00 2189-06-27 08:42:00
 4 WHITE                 NA                  NA
 5 WHITE                 2157-11-18 17:38:00 2157-11-19 01:24:00
 6 WHITE                 NA                  NA
 7 WHITE                 NA                  NA
 8 BLACK/AFRICAN AMERICAN 2131-01-07 13:36:00 2131-01-07 22:13:00
 9 OTHER                 NA                  NA
10 UNKNOWN               2162-02-17 19:35:00 2162-02-17 23:30:00
   hospital_expire_flag gender anchor_age anchor_year anchor_year_group
                  <int> <chr>      <int>       <int> <chr>
 1                    0 F             52        2180 2014 - 2016
 2                    0 F             86        2150 2008 - 2010
 3                    0 F             73        2186 2008 - 2010
 4                    0 F             55        2157 2011 - 2013
 5                    0 F             55        2157 2011 - 2013
 6                    0 F             46        2110 2011 - 2013
 7                    1 M             73        2131 2017 - 2019
 8                    1 F             68        2122 2008 - 2010
 9                    0 F             53        2156 2008 - 2010
10                    0 M             56        2162 2020 - 2022
   dod        hematocrit sodium bicarbonate chloride  wbc creatinine glucose
   <date>          <dbl>  <dbl>       <dbl>    <dbl> <dbl>      <dbl>   <dbl>
 1 2180-09-09       41.1    126          25       95  6.9        0.7     102
 2 2152-01-30       36.1    137          26      100  7.1        1        85
 3 2193-08-26       27.3    144          21      109  5.3        2.3      89
 4 NA               37.4    142          30      104  5.4        0.5      87
 5 NA               38.1    142          22      108 15.7        0.6     112
 6 NA                 NA    139          NA       98   NA         NA      NA
```

```
 7 2134-12-06       31.4    138           28         97  10.4          1.3        131
 8 2131-01-20       39.7    130           30         88  12.2          1.1        141
 9 NA               34.9    137           24        102   7.2          0.9        288
10 2162-12-11       34.3    125           18         NA  16.8          3.1         95
   potassium temperature_fahrenheit respiratory_rate
       <dbl>                  <dbl>            <dbl>
 1       6.7                   98.7             24
 2       4.8                   97.7             24.3
 3       3.9                   98               23.5
 4       4.1                   97.6             14
 5       4.2                   98.5             18
 6       4.1                   97.7             19
 7       3.9                   97.9             16.5
 8       4.5                   98.1             13
 9       3.5                   97.2             14
10       6.5                   97.9             21
   non_invasive_blood_pressure_diastolic non_invasive_blood_pressure_systolic
                                   <dbl>                                <dbl>
 1                                    48                                   84
 2                                    56.5                                106
 3                                   102                                  154
 4                                    93.3                                156
 5                                    90                                  151
 6                                    56                                   73
 7                                    78                                  110
 8                                    30.5                                174.
 9                                    62                                   98.5
10                                    80                                  112
   heart_rate age_intime
        <dbl>      <int>
 1         91         52
 2         78         86
 3         76         76
 4         79.3       55
 5         86         55
 6         86         46
 7        124.        76
 8         49         77
 9         80         57
10        110.        56
# i 94,448 more rows
```

Q1.8 Preprocessing

Perform the following preprocessing steps. (i) Lump infrequent levels into "Other" level for `first_careunit`, `last_careunit`, `admission_type`, `admission_location`, and `discharge_location`. (ii) Collapse the levels of `race` into `ASIAN`, `BLACK`, `HISPANIC`, `WHITE`, and `Other`. (iii) Create a new variable `los_long` that is `TRUE` when `los` is greater than or equal to 2 days. (iv) Summarize the data using `tbl_summary()`, stratified by `los_long`. Hint: `fct_lump_n` and `fct_collapse` from the `forcats` package are useful.

Hint: Below is a numerical summary of my tibble after preprocessing:

Solution

```
library('forcats')

mimic_icu_cohort |>
  #Lumping infrequent levels using fct_lump_n
  mutate(first_careunit = fct_lump_n(first_careunit, n = 4),
  last_careunit = fct_lump_n(last_careunit, n = 4),
  admission_type = fct_lump_n(admission_type, n = 4),
  admission_location = fct_lump_n(admission_location, n = 4),
  discharge_location = fct_lump_n(discharge_location, n = 4)) |>
  #Collapsing the levels of race
  mutate(race = fct_collapse(race,
  ASIAN = unique(mimic_icu_cohort$race)[grep('ASIAN',
  unique(mimic_icu_cohort$race))],
  BLACK = unique(mimic_icu_cohort$race)[grep('BLACK',
  unique(mimic_icu_cohort$race))],
  HISPANIC = unique(mimic_icu_cohort$race)[grep('HISPANIC',
  unique(mimic_icu_cohort$race))],
  WHITE = unique(mimic_icu_cohort$race)[grep('WHITE',
  unique(mimic_icu_cohort$race))],
  OTHER = unique(mimic_icu_cohort$race)[!grepl('ASIAN|BLACK|HISPANIC|WHITE',
  unique(mimic_icu_cohort$race))])) |>
  #Creating a variable los_long
  mutate(los_long = (los >= 2)) |>
  tbl_summary(by = los_long, include = c(first_careunit, last_careunit, los,
  admission_type, admission_location, discharge_location, insurance, language,
  marital_status, race, hospital_expire_flag, gender, dod, chloride, creatinine,
  sodium, potassium, glucose, hematocrit, wbc, bicarbonate,
  non_invasive_blood_pressure_systolic, non_invasive_blood_pressure_diastolic,
  respiratory_rate, temperature_fahrenheit, heart_rate, age_intime))
```

```
14 missing rows in the "los_long" column have been removed.
The following errors were returned during `tbl_summary()`:
x For variable `dod` (`los_long = FALSE`) and "p75" statistic: * not defined
  for "Date" objects
```

Q1.9 Save the final tibble

Save the final tibble to an R data file `mimic_icu_cohort.rds` in the `mimiciv_shiny` folder.

```
# make a directory mimiciv_shiny
if (!dir.exists("mimiciv_shiny")) {
  dir.create("mimiciv_shiny")
}
# save the final tibble
mimic_icu_cohort |>
  write_rds("mimiciv_shiny/mimic_icu_cohort.rds", compress = "gz")
```

Close database connection and clear workspace.

```
if (exists("con_bq")) {
  dbDisconnect(con_bq)
}
rm(list = ls())
```

Although it is not a good practice to add big data files to Git, for grading purpose, please add `mimic_icu_cohort.rds` to your Git repository.

Q2. Shiny app

Develop a Shiny app for exploring the ICU cohort data created in Q1. The app should reside in the `mimiciv_shiny` folder. The app should contain at least two tabs. One tab provides easy access to the graphical and numerical summaries of variables (demographics, lab measurements, vitals) in the ICU cohort, using the `mimic_icu_cohort.rds` you curated in Q1. The other tab allows user to choose a specific patient in the cohort and display the patient's ADT and ICU stay information as we did in Q1 of HW3, by dynamically retrieving the patient's ADT and ICU stay information from BigQuery database. Again, do **not** ever add the BigQuery token to your Git repository. If you do so, you will lose 50 points.

19

Characteristic	**TRUE** N = 46,337[1]	F
first_careunit		
Cardiac Vascular Intensive Care Unit (CVICU)	7,353 (16%)	
Medical Intensive Care Unit (MICU)	9,837 (21%)	
Medical/Surgical Intensive Care Unit (MICU/SICU)	6,667 (14%)	
Surgical Intensive Care Unit (SICU)	6,434 (14%)	
Other	16,046 (35%)	
last_careunit		
Cardiac Vascular Intensive Care Unit (CVICU)	7,353 (16%)	
Medical Intensive Care Unit (MICU)	9,837 (21%)	
Medical/Surgical Intensive Care Unit (MICU/SICU)	6,667 (14%)	
Surgical Intensive Care Unit (SICU)	6,434 (14%)	
Other	16,046 (35%)	
los	3.9 (2.7, 6.8)	
admission_type		
EW EMER.	23,012 (50%)	
OBSERVATION ADMIT	7,393 (16%)	
SURGICAL SAME DAY ADMISSION	4,001 (8.6%)	
URGENT	8,691 (19%)	
Other	3,240 (7.0%)	
admission_location		
EMERGENCY ROOM	17,058 (37%)	
PHYSICIAN REFERRAL	11,013 (24%)	
TRANSFER FROM HOSPITAL	13,904 (30%)	
WALK-IN/SELF REFERRAL	2,169 (4.7%)	
Other	2,193 (4.7%)	
discharge_location		
DIED	6,884 (15%)	
HOME	6,879 (15%)	
HOME HEALTH CARE	10,620 (23%)	
SKILLED NURSING FACILITY	8,785 (19%)	
Other	13,092 (28%)	
Unknown	77	
insurance		
Medicaid	6,768 (15%)	
Medicare	26,330 (58%)	
No charge	5 (<0.1%)	
Other	1,091 (2.4%)	
Private	11,515 (25%)	
Unknown	628	
language		
American Sign Language	29 (<0.1%)	
Amharic	14 (<0.1%)	
Arabic	87 (0.2%)	
Armenian	12 (<0.1%)	
Bengali	22 (<0.1%)	
Chinese	550 (1.2%)	
English	41,563 (90%)	
French	18 (<0.1%)	
Haitian	375 (0.8%)	